

Ouro: Personalized Chatbot with Dynamic Responses

Final Technical Project Report

Project ID: Generative AI - Project 2

Team Lead: Pierre Ramez

Date: November 27, 2025

Status: Production Ready

Executive Summary

The **Ouro Project** was engineered to address the limitations of static Large Language Models (LLMs) in personalized environments. While traditional chatbots require expensive, periodic retraining to update knowledge or correct behaviors, Ouro implements a **Closed-Loop Active Learning System**. This system leverages a decoupled microservices architecture to enable a Human-in-the-Loop (HITL) feedback mechanism. By utilizing Low-Rank Adaptation (LoRA) and 4-bit Quantization via the Unsloth framework, we achieved a high-performance conversational agent (Llama 3.2 3B) capable of running on resource-constrained edge infrastructure (CPU/T4 GPU) while maintaining the ability to “heal” its own hallucinations through real-time user corrections.

Milestone 1: Data Collection and Preprocessing

Objective: To curate a high-quality, heterogeneous dataset that balances general conversational fluency with specific emotional intelligence capabilities.

1.1 Data Sourcing Strategy

We aggregated data from three distinct sources to create a robust training corpus:

- **Daily Dialogue (CSV):** A large-scale dataset (13,118 rows) containing multi-turn conversations about daily life, providing the “chitchat” backbone.
- **Mental Health Dataset (JSON):** 661 intent-response pairs focused on emotional support and empathy.
- **Human Chat (TXT):** 94 raw logs of human-to-human interaction to capture realistic turn-taking dynamics.

1.2 The Preprocessing Pipeline

A custom ETL (Extract, Transform, Load) pipeline was engineered using pandas and ftfy to sanitize the raw inputs. Key Transformation Steps:

- **Encoding Normalization:** Utilized ftfy to fix “mojibake” (broken unicode characters) and standardize text encoding.
- **Noise Reduction:** Removed dataset-specific metadata columns (e.g., act, emotion from Daily Dialogue) to focus purely on semantic content.
- **Regex Cleaning:**
 - Collapsed extraneous whitespace.
 - Standardized punctuation (e.g., converting “Hello .” → “Hello.”).
 - Removed artifact quotes surrounding dialogue turns.

1.3 Final Dataset Schema

The heterogeneous sources were unified into the industry-standard Llama-3 Chat Format:

- **Total Volume:** 13,502 high-quality conversational pairs.
- **Format:** JSONL (JSON Lines).
- **Structure:**

```
{
  "messages": [
    { "role": "user", "content": "I feel overwhelmed today." },
    { "role": "assistant", "content": "I understand. Let's take it one step at a time." }
  ]
}
```

Milestone 2: Model Development and Training

Objective: To fine-tune a Transformer-based model that balances reasoning capability with inference latency constraints.

2.1 Comparative Architecture Analysis

We conducted a rigorous *A/B* test between two architectures to determine the optimal production model.

Candidate A: Llama 3.1 8B (The “Reasoning” Model)

- **Configuration:** LoRA Rank 16, Alpha 16, Flash Attention 2.
- **Performance:** High reasoning capability.
- **Constraint:** Required ~14GB VRAM for inference.
- **Outcome:** Rejected for production due to memory constraints on the target deployment environment (Hugging Face Spaces Free Tier supports only 16GB CPU RAM).

Candidate B: Llama 3.2 3B (The “Edge” Model)

- **Configuration:** 4-bit NF4 Quantization, LoRA Rank 16.
- **Performance:** Excellent conversational coherence with significantly lower footprint.
- **Constraint:** Fits comfortably within 16GB CPU RAM.
- **Outcome:** Selected for production.

2.2 The Training Pipeline (Unsloth Optimization)

We utilized *Unsloth*, a framework that manually optimizes backpropagation steps, allowing for 2x faster training and 60% memory reduction compared to standard Hugging Face implementations.

Hyperparameter Configuration:

Parameter	Value	Justification
Precision	bf16 (Mixed)	Balances numerical stability with memory usage.
Quantization	4-bit (NF4)	Reduces VRAM usage from ~6GB to ~2.5GB.
LoRA Rank (r)	16	Sufficient plasticity for style transfer without overfitting.
Optimizer	adamw_8bit	Further reduces optimizer state memory footprint.
Max Seq Length	2048	Covers standard chat history context windows.

2.3 Evaluation Strategy: The Sliding Window

We rejected standard “final-turn” evaluation in favor of a Sliding Window Strategy. This involves taking a multi-turn conversation and evaluating the model’s prediction accuracy at every assistant turn, not just the end.

Quantitative Results:

- **Perplexity:** 4.08 (Indicates high predictive confidence).
- **BERTScore F1:** 0.8769 (Indicates ~ 88% semantic alignment with human references).
- **Precision:** 0.8849 | **Recall:** 0.8693.

Milestone 3: Advanced Techniques (Continuous Learning)

Objective: To architect a “Self-Healing” AI pipeline that integrates user corrections. This milestone represents the project’s core innovation. Instead of a static model, Ouro implements a dynamic feedback loop.

3.1 The Feedback Architecture

- **Interaction:** User chats with the bot via the Frontend.
- **Intervention:** The user identifies an error (hallucination) and submits a “Correction” via the UI.
- **Ingestion:** The backend captures this correction into a persistent feedback.jsonl ledger.
- **Orchestration:** A custom PipelineRun class triggers the learning workflow:
 - **Data Preparation:** Converts the raw correction into a training batch.
 - **Active Training:** Initiates a targeted SFT run on the specific correction data.
 - **Validation:** Checks if the new adapter passes “Quality Gates” (Semantic Similarity > 0.85).
 - **Hot-Swap:** The system unloads the old LoRA adapter and loads the new one without restarting the server.

3.2 Dynamic Adapter Management

By leveraging PeftModel, we decoupled the base model weights (frozen) from the behavior weights (adapters). This allows us to “patch” the model’s behavior in seconds rather than hours.

Milestone 4: MLOps and Model Management

Objective: To ensure reproducibility, traceability, and safe deployment.

4.1 Experiment Tracking (MLflow)

We integrated MLflow (hosted via DagsHub) to track every training run.

- **Metrics Logged:** Loss curves, Learning Rate, ROUGE-L, BLEU, BERTScore.
- **Artifacts:** The trained adapter files and tokenizer configurations are versioned and stored.
- **Registry:** Only models that pass the automated Quality Gates are tagged as production_candidate.

4.2 CI/CD Pipeline (GitHub Actions)

We implemented a robust Continuous Integration/Continuous Deployment pipeline defined in `main.yml`.

Workflow:

- **Trigger:** Code push to main branch.
- **Build & Test:**
 - Sets up Python environment.
 - Installs dependencies.
 - Runs `pytest` on `test_app.py` to verify API endpoints (`/health`, `/chat`, `/feedback`).
- **Deploy:**
 - If tests pass, the code is pushed to the Hugging Face Space.
 - The Docker container is automatically rebuilt and redeployed.

4.3 Containerization

The application is packaged using Docker to ensure consistency across environments (Colab, Local, Cloud).

- **Permissions:** Configured `chmod 777` on the `/data` directory to ensure the application can write feedback logs even in restricted serverless environments.

Milestone 5: Final Deployment

Objective: To deliver a user-friendly, accessible product.

5.1 Deployment Stack

- **Frontend:** Hugging Face Spaces (Streamlit). Provides a chat interface with a custom “Spotted a mistake?” expander.
- **Backend:** Hugging Face Spaces (FastAPI). Handles inference, queueing, and feedback logic.
- **Model Registry:** Hugging Face Hub. Stores the Llama 3.2 3B base model and the fine-tuned LoRA adapters.

5.2 Zero-Downtime Updates

The deployment architecture supports Hot Reloading. When a new adapter is trained:

1. The pipeline pushes the adapter to the Hub.
2. The backend receives a `POST /reload-adapter` request.
3. The inference engine swaps the weights in memory immediately.
4. The user experiences improved performance instantly.

5.3 Links & Resources

- **GitHub Repo:** [PierreRamez/PersonalChatbot](#)
- **Production Model:** [pierrramez/Llama-3.2-3B-Instruct-bnb-4bit_finetuned](#)
- **Live Demo:** [Chatbot UI](#)
- **MLflow Tracking:** [DagsHub Project](#)

Conclusion

The Ouro Project successfully transcends the typical limitations of academic AI projects. By integrating a Continuous Learning Loop, we have moved beyond creating a “smart” chatbot to creating a “learning” chatbot. The combination of Unsloth for efficient training, LoRA for modular updates, and a CI/CD pipeline for safe deployment results in a resilient conversational agent that adapts and improves with real-user feedback and a measured quality-gate pipeline. Ouro demonstrates that with careful adapter design and automated validation, a production-grade personalization loop is feasible on resource-constrained infrastructure.

Future Work

Latency & Serving Optimization

- Move inference from CPU-only Hugging Face Spaces to a T4/GPU-backed deployment (or scaled CPU cluster) to reduce per-response latency from ~1 minute to sub-5s.
- Explore batching, request-level caching, and async streaming responses.
- Investigate alternative quantization schemes (QLoRA workflows, GPTQ) and kernel-level optimizations to further reduce latency.

Robustness & Safety

- Implement adversarial testing and synthetic negation cases to harden hallucination detection.
- Add content filters and safety classifiers upstream of the model to avoid undesirable outputs.

Evaluation & User Studies

- Run longitudinal user studies to quantify improvement rates from the active feedback loop.
- Measure user trust and perceived usefulness across different demographics.

Privacy & Compliance

- Add differential privacy mechanisms to the feedback ingestion pipeline to protect personally identifiable information in user corrections.
- Formalize data retention policies and consent flows for production users.

Feature Roadmap

- Multimodal inputs (image + text) for richer conversational grounding.
- Transfer learning across adapters to allow knowledge sharing between domain- specific agents.
- Plugin-style external tool invocation (calendars, search, calculators).

Acknowledgements

Thanks to the DEPI GenAI diploma program, our university mentors, and all the team members for their time and feedback during development.

How to reproduce

- Clone the repository: `git clone https://github.com/PierreRamez/PersonalChatbot`
- Install dependencies: `pip install -r requirements.txt`
- Prepare datasets: follow docs/data_prep.md
- Run training: `python train_adapter.py --config configs/unsloth_lora_3b.json`
- Deploy: follow docs/deploy_hf_space.md for Hugging Face Spaces CPU deployment.

Contact

Pierre Ramez — pierre.ramez10@gmail.com